

Practical Machine Learning - Course Project

Alasdair Hazen

Executive Summary

Six subjects participated in a study of exercising with dumbbells. The experiment examines, not the quantity, but the manner in which they performed the exercise. There were five manners in which the subjects could do the exercise. (A) exactly to the specification, (B) throwing the elbows to the front, (C) lifting the dumbbell only halfway, (D) lowering the dumbbell only halfway, and (E) throwing the hips to the front. The goal of this project is to analyse the data and predict which manner they performed the exercise based on the data. For this we will construct several models and select the best one according to accuracy.

Question

By analysing the data from accelerometers on the belt, forearm, arm, and dumbbell using an algorithm, can the appropriate activity quality class be predicted?

Input Data

1. Set libraries
 - caret
 - rattle
 - rpart
 - randomForest
 - class
2. Download data
3. Read data into datasets
4. Drop unnecessary columns
 - NAs
 - blanks
 - errors
5. Split testing data into smaller group(s) to build a model and test the model

```
# Check packages are installed
# any(grepl("caret", installed.packages()))
# any(grepl("rattle", installed.packages()))
# any(grepl("rpart", installed.packages()))
# any(grepl("rpart.plot", installed.packages()))
# any(grepl("randomForest", installed.packages()))

# Set libraries
library(caret)
library(rattle)
library(rpart)
library(rpart.plot)
library(randomForest)
library(class)
```

```
# Download data
if (!file.exists("pml-training.csv")) {
  download.file("http://d396qusza40orc.cloudfront.net/predmachlearn/pml-trainin
g.csv",
               destfile = "pml-training.csv")
}

if (!file.exists("pml-testing.csv")) {
  download.file("http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testin
g.csv",
               destfile = "pml-testing.csv")
}
```

```
# Read data into datasets
pmlTrainingData <- read.csv("pml-training.csv", na.strings=c("NA",""), header=T
RUE)
pmlTestingData <- read.csv("pml-testing.csv", na.strings=c("NA",""), header=TRU
E)
```

```
# Drop unnecessary columns
pmlTrainingData <- pmlTrainingData[,!(names(pmlTrainingData) %in% drop)]
pmlTrainingData <- pmlTrainingData[,8:length(colnames(pmlTrainingData))]
```

```
pmlTestingData <- pmlTestingData[,!(names(pmlTestingData) %in% drop)]
pmlTestingData <- pmlTestingData[,8:length(colnames(pmlTestingData))]
```

We will do a quick check that our testing data and our training data columns are the same.

```
# Show remaining columns.
colnames(pmlTrainingData)
```

```
## [1] "roll_belt"           "pitch_belt"           "yaw_belt"
## [4] "total_accel_belt"    "gyros_belt_x"         "gyros_belt_y"
## [7] "gyros_belt_z"        "accel_belt_x"         "accel_belt_y"
## [10] "accel_belt_z"        "magnet_belt_x"        "magnet_belt_y"
## [13] "magnet_belt_z"       "roll_arm"             "pitch_arm"
## [16] "yaw_arm"             "total_accel_arm"      "gyros_arm_x"
## [19] "gyros_arm_y"         "gyros_arm_z"          "accel_arm_x"
## [22] "accel_arm_y"         "accel_arm_z"          "magnet_arm_x"
## [25] "magnet_arm_y"        "magnet_arm_z"         "roll_dumbbell"
## [28] "pitch_dumbbell"      "yaw_dumbbell"         "total_accel_dumbbell"
## [31] "gyros_dumbbell_x"    "gyros_dumbbell_y"     "gyros_dumbbell_z"
## [34] "accel_dumbbell_x"    "accel_dumbbell_y"     "accel_dumbbell_z"
## [37] "magnet_dumbbell_x"   "magnet_dumbbell_y"    "magnet_dumbbell_z"
## [40] "roll_forearm"        "pitch_forearm"        "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x"      "gyros_forearm_y"
## [46] "gyros_forearm_z"     "accel_forearm_x"      "accel_forearm_y"
## [49] "accel_forearm_z"     "magnet_forearm_x"     "magnet_forearm_y"
## [52] "magnet_forearm_z"    "classe"
```

```
colnames(pmlTestingData)
```

```
## [1] "roll_belt"           "pitch_belt"           "yaw_belt"
## [4] "total_accel_belt"    "gyros_belt_x"         "gyros_belt_y"
## [7] "gyros_belt_z"        "accel_belt_x"         "accel_belt_y"
## [10] "accel_belt_z"        "magnet_belt_x"        "magnet_belt_y"
## [13] "magnet_belt_z"       "roll_arm"             "pitch_arm"
## [16] "yaw_arm"             "total_accel_arm"      "gyros_arm_x"
## [19] "gyros_arm_y"         "gyros_arm_z"          "accel_arm_x"
## [22] "accel_arm_y"         "accel_arm_z"          "magnet_arm_x"
## [25] "magnet_arm_y"        "magnet_arm_z"         "roll_dumbbell"
## [28] "pitch_dumbbell"      "yaw_dumbbell"         "total_accel_dumbbell"
## [31] "gyros_dumbbell_x"    "gyros_dumbbell_y"     "gyros_dumbbell_z"
## [34] "accel_dumbbell_x"    "accel_dumbbell_y"     "accel_dumbbell_z"
## [37] "magnet_dumbbell_x"   "magnet_dumbbell_y"    "magnet_dumbbell_z"
## [40] "roll_forearm"        "pitch_forearm"        "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x"      "gyros_forearm_y"
## [46] "gyros_forearm_z"     "accel_forearm_x"      "accel_forearm_y"
## [49] "accel_forearm_z"     "magnet_forearm_x"     "magnet_forearm_y"
## [52] "magnet_forearm_z"    "problem_id"
```

```
# show classe
table(pmlTrainingData$classe)
```

```
##  
##      A      B      C      D      E  
## 5580 3797 3422 3216 3607
```

```
# show probability  
round(prop.table(table(pmlTrainingData$classe)) * 100, digits = 1)
```

```
##  
##      A      B      C      D      E  
## 28.4 19.4 17.4 16.4 18.4
```

Here we can see that the last column differs. In Training data the last column is `classe`, whereas in Testing data the column is `problem_id`. We will need to adjust for this.

There are several ways we can partition the data. We can use folds, partitions, etc. and a variety of variations and combinations. However, after 16 tries I have decided that the simplest 60/40 partition is just as effective as a more complicated derivative.

```
# Divide each of these 4 sets into training (60%) and test (40%) sets.  
set.seed(1979)  
inTrain <- createDataPartition(y=pmlTrainingData$classe, p=0.6, list=FALSE)  
training <- pmlTrainingData[inTrain,]  
testing <- pmlTrainingData[-inTrain,]
```

Features

1. Convert/Create covariates
 - Covariates that have no variability (NearZeroVar)
2. Check for overfitting

```
# Check NearZeroVar  
nzv <- nearZeroVar(pmlTrainingData, saveMetrics=TRUE)  
nzv
```

##	freqRatio	percentUnique	zeroVar	nzv
## roll_belt	1.101904	6.7781062	FALSE	FALSE
## pitch_belt	1.036082	9.3772296	FALSE	FALSE
## yaw_belt	1.058480	9.9734991	FALSE	FALSE
## total_accel_belt	1.063160	0.1477933	FALSE	FALSE
## gyros_belt_x	1.058651	0.7134849	FALSE	FALSE
## gyros_belt_y	1.144000	0.3516461	FALSE	FALSE
## gyros_belt_z	1.066214	0.8612782	FALSE	FALSE
## accel_belt_x	1.055412	0.8357966	FALSE	FALSE
## accel_belt_y	1.113725	0.7287738	FALSE	FALSE
## accel_belt_z	1.078767	1.5237998	FALSE	FALSE
## magnet_belt_x	1.090141	1.6664968	FALSE	FALSE
## magnet_belt_y	1.099688	1.5187035	FALSE	FALSE
## magnet_belt_z	1.006369	2.3290184	FALSE	FALSE
## roll_arm	52.338462	13.5256345	FALSE	FALSE
## pitch_arm	87.256410	15.7323412	FALSE	FALSE
## yaw_arm	33.029126	14.6570176	FALSE	FALSE
## total_accel_arm	1.024526	0.3363572	FALSE	FALSE
## gyros_arm_x	1.015504	3.2769341	FALSE	FALSE
## gyros_arm_y	1.454369	1.9162165	FALSE	FALSE
## gyros_arm_z	1.110687	1.2638875	FALSE	FALSE
## accel_arm_x	1.017341	3.9598410	FALSE	FALSE
## accel_arm_y	1.140187	2.7367241	FALSE	FALSE
## accel_arm_z	1.128000	4.0362858	FALSE	FALSE
## magnet_arm_x	1.000000	6.8239731	FALSE	FALSE
## magnet_arm_y	1.056818	4.4439914	FALSE	FALSE
## magnet_arm_z	1.036364	6.4468454	FALSE	FALSE
## roll_dumbbell	1.022388	84.2065029	FALSE	FALSE
## pitch_dumbbell	2.277372	81.7449801	FALSE	FALSE
## yaw_dumbbell	1.132231	83.4828254	FALSE	FALSE
## total_accel_dumbbell	1.072634	0.2191418	FALSE	FALSE
## gyros_dumbbell_x	1.003268	1.2282132	FALSE	FALSE
## gyros_dumbbell_y	1.264957	1.4167771	FALSE	FALSE
## gyros_dumbbell_z	1.060100	1.0498420	FALSE	FALSE
## accel_dumbbell_x	1.018018	2.1659362	FALSE	FALSE
## accel_dumbbell_y	1.053061	2.3748853	FALSE	FALSE
## accel_dumbbell_z	1.133333	2.0894914	FALSE	FALSE
## magnet_dumbbell_x	1.098266	5.7486495	FALSE	FALSE
## magnet_dumbbell_y	1.197740	4.3012945	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.4451126	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE
## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE
## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE

```
## accel_forearm_y      1.059406      5.1116094      FALSE FALSE
## accel_forearm_z      1.006250      2.9558659      FALSE FALSE
## magnet_forearm_x      1.012346      7.7667924      FALSE FALSE
## magnet_forearm_y      1.246914      9.5403119      FALSE FALSE
## magnet_forearm_z      1.000000      8.5771073      FALSE FALSE
## classe                1.469581      0.0254816      FALSE FALSE
```

Near zero variance showing all false shows our data is clean after we removed superfluous columns.
No more cleaning is required.

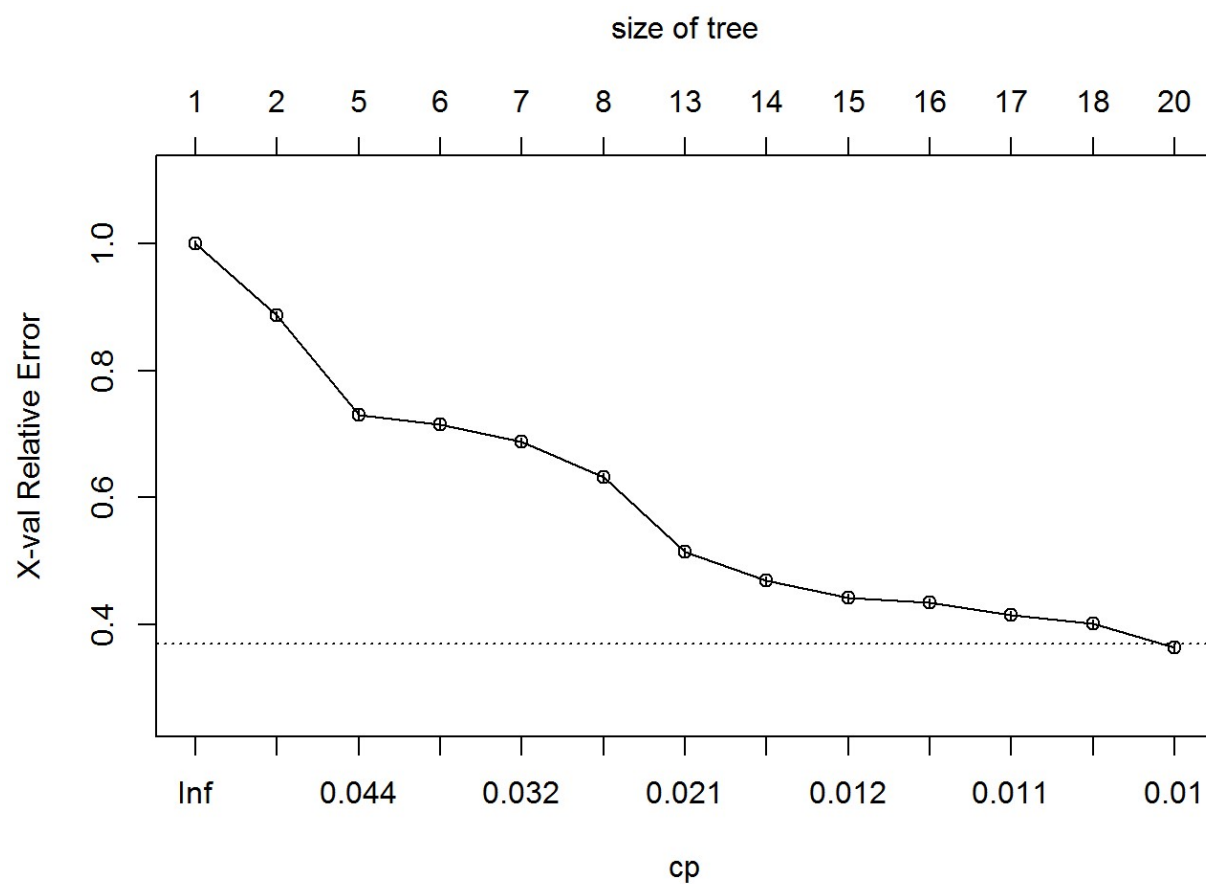
Algorithm

1. Cart Modeling via rpart
2. classification tree with (method=rpart)
3. Regression
4. Random forest (method=rf)

```
# grow tree
fit <- rpart(classe ~ .,method="class",data=training)
printcp(fit) # display the results
```

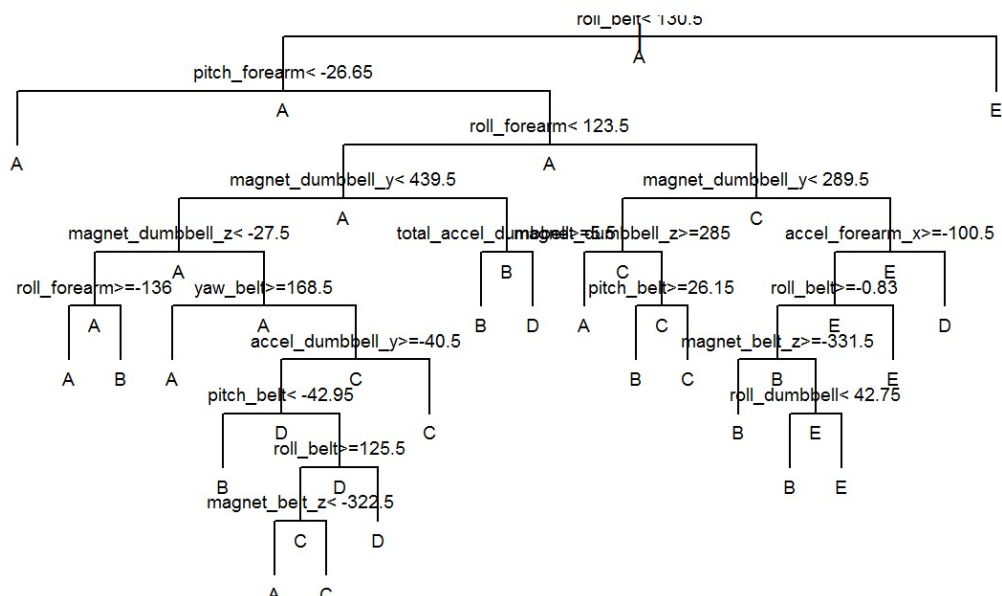
```
##
## Classification tree:
## rpart(formula = classe ~ ., data = training, method = "class")
##
## Variables actually used in tree construction:
## [1] accel_dumbbell_y      accel_forearm_x      magnet_belt_z
## [4] magnet_dumbbell_y     magnet_dumbbell_z    pitch_belt
## [7] pitch_forearm         roll_belt             roll_dumbbell
## [10] roll_forearm          total_accel_dumbbell yaw_belt
##
## Root node error: 8428/11776 = 0.71569
##
## n= 11776
##
##          CP nsplit rel error  xerror      xstd
## 1  0.113075      0  1.00000 1.00000 0.0058081
## 2  0.050724      1  0.88692 0.88728 0.0061987
## 3  0.038087      4  0.72057 0.73113 0.0064309
## 4  0.034053      5  0.68249 0.71583 0.0064359
## 5  0.030612      6  0.64843 0.68747 0.0064371
## 6  0.022781      7  0.61782 0.63253 0.0064090
## 7  0.018747     12  0.49573 0.51448 0.0062102
## 8  0.012696     13  0.47698 0.46868 0.0060792
## 9  0.011984     14  0.46429 0.44198 0.0059878
## 10 0.011391     15  0.45230 0.43439 0.0059596
## 11 0.010916     16  0.44091 0.41445 0.0058813
## 12 0.010679     17  0.43000 0.40081 0.0058236
## 13 0.010000     19  0.40864 0.36308 0.0056467
```

```
# plot tree
plotcp(fit) # visualize cross-validation results
```



```
# summary(fit) # detailed summary of splits (For testing only. Lengthy output)
plot(fit, uniform=TRUE,
      main="Classification Tree")
text(fit, use.n=FALSE, all=TRUE, cex=.6)
```


Classification Tree

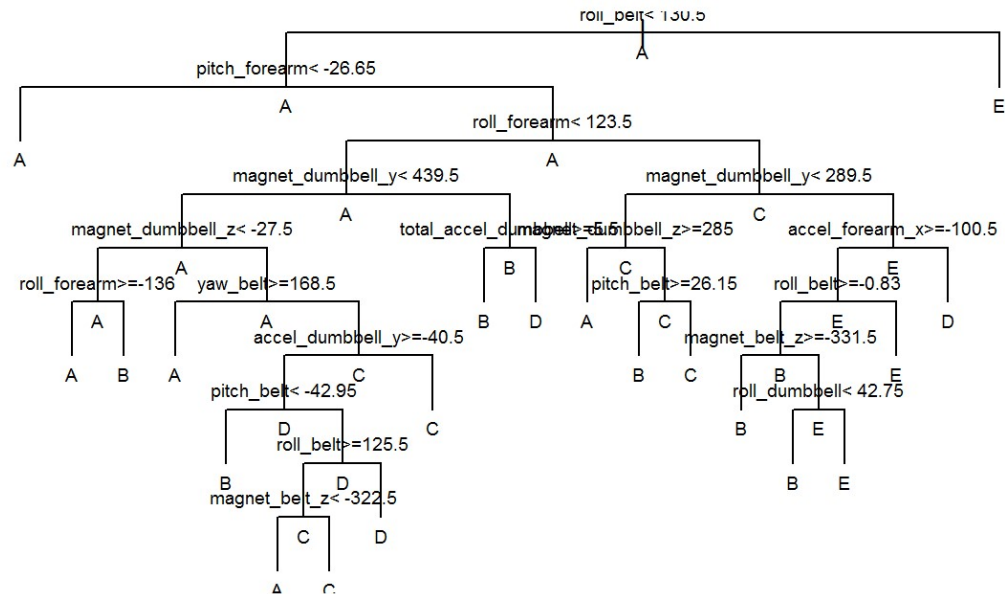


```

# prune the tree
pfit<- prune(fit, cp= fit$testing[which.min(fit$testing[, "xerror"]), "CP"])
# plot the pruned tree
plot(pfit, uniform=TRUE,
     main="Pruned Classification Tree")
text(pfit, use.n=FALSE, all=TRUE, cex=.6)

```

Pruned Classification Tree



Parameters

1. cross-validation
2. bootstrapping

```
# Random Forest prediction of training data
fit <- randomForest(classe ~ ., data=training)
print(fit) # view results
```

```
##
## Call:
##  randomForest(formula = classe ~ ., data = training)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 7
##
##              OOB estimate of  error rate: 0.65%
## Confusion matrix:
```

	A	B	C	D	E	class.error
A	3345	3	0	0	0	0.0008960573
B	16	2258	5	0	0	0.0092145678
C	0	16	2032	6	0	0.0107108082
D	0	0	20	1909	1	0.0108808290
E	0	0	0	10	2155	0.0046189376

```
importance(fit) # importance of each predictor
```

##	MeanDecreaseGini
## roll_belt	754.89965
## pitch_belt	406.49451
## yaw_belt	520.85462
## total_accel_belt	119.80911
## gyros_belt_x	61.53216
## gyros_belt_y	66.83782
## gyros_belt_z	180.41815
## accel_belt_x	75.61335
## accel_belt_y	74.75943
## accel_belt_z	249.44779
## magnet_belt_x	151.12180
## magnet_belt_y	230.18909
## magnet_belt_z	227.10018
## roll_arm	182.20775
## pitch_arm	99.67084
## yaw_arm	147.51201
## total_accel_arm	59.51907
## gyros_arm_x	77.71598
## gyros_arm_y	81.91224
## gyros_arm_z	38.71183
## accel_arm_x	137.14882
## accel_arm_y	93.25851
## accel_arm_z	75.99641
## magnet_arm_x	147.90905
## magnet_arm_y	128.33402
## magnet_arm_z	115.48929
## roll_dumbbell	251.42958
## pitch_dumbbell	105.70308
## yaw_dumbbell	164.72312
## total_accel_dumbbell	148.47264
## gyros_dumbbell_x	80.35083
## gyros_dumbbell_y	150.19494
## gyros_dumbbell_z	52.67296
## accel_dumbbell_x	150.41538
## accel_dumbbell_y	242.86676
## accel_dumbbell_z	194.76933
## magnet_dumbbell_x	297.72593
## magnet_dumbbell_y	399.45126
## magnet_dumbbell_z	445.67412
## roll_forearm	383.99512
## pitch_forearm	475.21346
## yaw_forearm	111.62573
## total_accel_forearm	67.13098
## gyros_forearm_x	48.07885
## gyros_forearm_y	79.73248
## gyros_forearm_z	50.33325
## accel_forearm_x	198.90877

```
## accel_forearm_y          94.57658
## accel_forearm_z         151.64496
## magnet_forearm_x        138.50751
## magnet_forearm_y        136.13522
## magnet_forearm_z        184.30248
```

Train Model

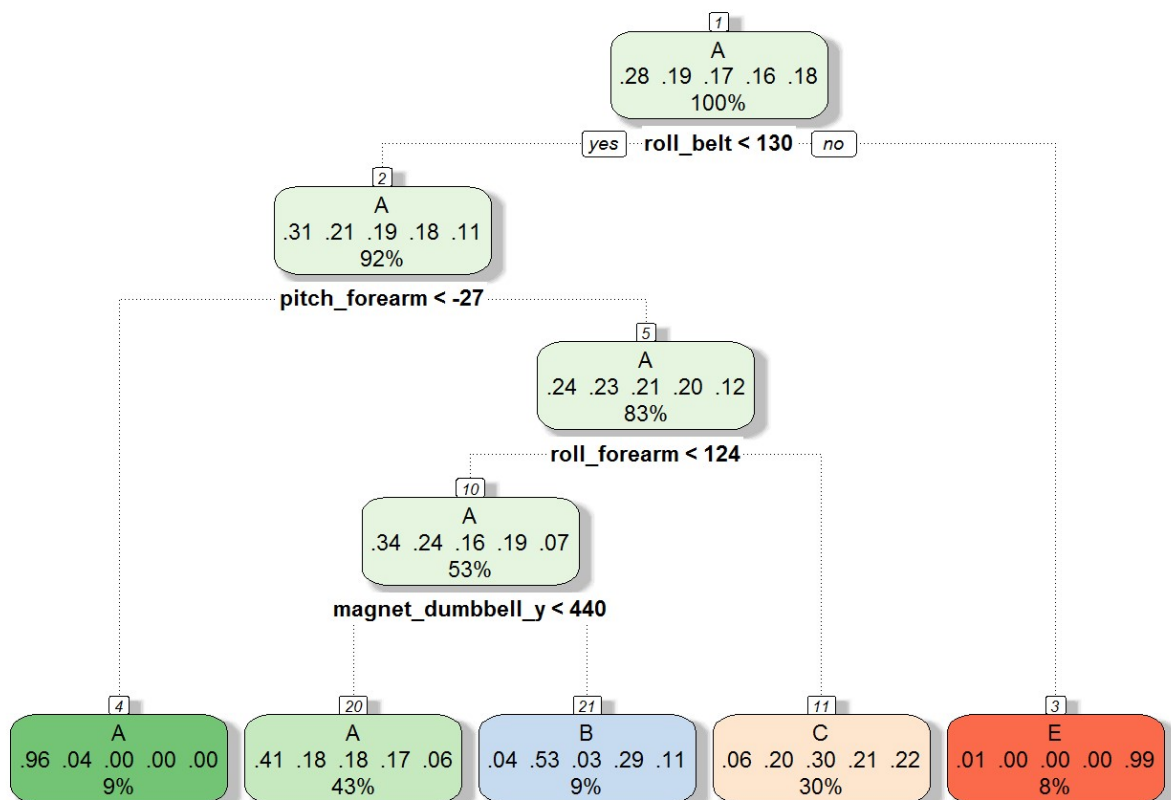
```
# Run rpart
set.seed(1979)
testModel <- train(training$classe ~ ., data = training, method="rpart")
print(testModel, digits=3)
```

```
## CART
##
## 11776 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 11776, 11776, 11776, 11776, 11776, 11776, ...
## Resampling results across tuning parameters:
##
##    cp      Accuracy  Kappa  Accuracy SD  Kappa SD
##    0.0381  0.491     0.3334  0.0405      0.0658
##    0.0507  0.443     0.2558  0.0607      0.1028
##    0.1131  0.330     0.0711  0.0395      0.0593
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0381.
```

```
print(testModel$finalModel, digits=3)
```

```
## n= 11776
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 11776 8430 A (0.28 0.19 0.17 0.16 0.18)
##    2) roll_belt< 130 10803 7460 A (0.31 0.21 0.19 0.18 0.11)
##      4) pitch_forearm< -26.6 1030 46 A (0.96 0.045 0 0 0) *
##      5) pitch_forearm>=-26.6 9773 7420 A (0.24 0.23 0.21 0.2 0.12)
##        10) roll_forearm< 124 6228 4100 A (0.34 0.24 0.16 0.19 0.069)
##          20) magnet_dumbbell_y< 440 5119 3030 A (0.41 0.18 0.18 0.17 0.061) *
##          21) magnet_dumbbell_y>=440 1109 518 B (0.04 0.53 0.029 0.29 0.11) *
##        11) roll_forearm>=124 3545 2460 C (0.063 0.2 0.3 0.21 0.22) *
##    3) roll_belt>=130 973 10 E (0.01 0 0 0 0.99) *
```

```
fancyRpartPlot(testModel$finalModel)
```



Rattle 2016-Jan-27 17:20:59 alasd

```
# Run with no extra features.
predictions <- predict(testModel, newdata=testing)
print(confusionMatrix(predictions, testing$classe), digits=4)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A     B     C     D     E
##           A 2032  679  645  591  213
##           B   36  359   21  218   91
##           C  160  480  702  477  470
##           D    0    0    0    0    0
##           E    4    0    0    0  668
##
## Overall Statistics
##
##           Accuracy : 0.4794
##           95% CI : (0.4682, 0.4905)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.3191
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9104  0.23650  0.51316  0.0000  0.46325
## Specificity          0.6209  0.94216  0.75502  1.0000  0.99938
## Pos Pred Value       0.4885  0.49517  0.30668      NaN  0.99405
## Neg Pred Value       0.9457  0.83724  0.88015  0.8361  0.89211
## Prevalence           0.2845  0.19347  0.17436  0.1639  0.18379
## Detection Rate       0.2590  0.04576  0.08947  0.0000  0.08514
## Detection Prevalence 0.5302  0.09240  0.29174  0.0000  0.08565
## Balanced Accuracy     0.7657  0.58933  0.63409  0.5000  0.73131
```

```
# Train with only preprocessing.
set.seed(1979)
testModell <- train(training$classe ~ ., preProcess=c("center", "scale"), dat
a = training, method="rpart")
print(testModell, digits=3)
```

```
## CART
##
## 11776 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered (52), scaled (52)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 11776, 11776, 11776, 11776, 11776, 11776, ...
## Resampling results across tuning parameters:
##
##    cp      Accuracy  Kappa  Accuracy SD  Kappa SD
##    0.0381  0.491     0.3334  0.0405      0.0658
##    0.0507  0.443     0.2558  0.0607      0.1028
##    0.1131  0.330     0.0711  0.0395      0.0593
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was cp = 0.0381.
```

```
# Train with only cross validation.
set.seed(1979)
testModel2 <- train(training$classe ~ ., trControl=trainControl(method = "c
v", number = 4), data = training, method="rpart")
print(testModel, digits=3)
```

```
## CART
##
## 11776 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 11776, 11776, 11776, 11776, 11776, 11776, ...
## Resampling results across tuning parameters:
##
##    cp      Accuracy  Kappa  Accuracy SD  Kappa SD
##    0.0381  0.491     0.3334  0.0405      0.0658
##    0.0507  0.443     0.2558  0.0607      0.1028
##    0.1131  0.330     0.0711  0.0395      0.0593
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was cp = 0.0381.
```



```
# Train with both preprocessing and cross validation.
set.seed(1979)
testModel2 <- train(training$classe ~ ., preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number = 4), data = training, method="rpart")
print(testModel, digits=3)
```

```
## CART
##
## 11776 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 11776, 11776, 11776, 11776, 11776, 11776, ...
## Resampling results across tuning parameters:
##
##    cp          Accuracy  Kappa    Accuracy SD   Kappa SD
##    0.0381  0.491        0.3334  0.0405        0.0658
##    0.0507  0.443        0.2558  0.0607        0.1028
##    0.1131  0.330        0.0711  0.0395        0.0593
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0381.
```

```
# Train with only cross validation.
set.seed(1979)
testModel3 <- train(training$classe ~ ., method="rf", trControl=trainControl(method = "cv", number = 4), data=training)
print(testModel3, digits=3)
```

```
## Random Forest
##
## 11776 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 8832, 8831, 8831, 8834
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa  Accuracy SD  Kappa SD
##    2    0.988    0.985   0.00210    0.00265
##   27    0.987    0.984   0.00154    0.00195
##   52    0.979    0.974   0.00358    0.00453
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 2.
```

Evaluation

Now that we've made our model, we put our subset test data into model.

```
# Run against testModel3.
predictions <- predict(testModel3, newdata=testing)
print(confusionMatrix(predictions, testing$classe), digits=4)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A     B     C     D     E
##           A 2231    17     0     0     0
##           B   0 1498    10     0     0
##           C   0     3 1355    21     1
##           D   0     0     3 1265     6
##           E   1     0     0     0 1435
##
## Overall Statistics
##
##           Accuracy : 0.9921
##           95% CI : (0.9899, 0.9939)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.99
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9996   0.9868   0.9905   0.9837   0.9951
## Specificity          0.9970   0.9984   0.9961   0.9986   0.9998
## Pos Pred Value       0.9924   0.9934   0.9819   0.9929   0.9993
## Neg Pred Value       0.9998   0.9968   0.9980   0.9968   0.9989
## Prevalence           0.2845   0.1935   0.1744   0.1639   0.1838
## Detection Rate       0.2843   0.1909   0.1727   0.1612   0.1829
## Detection Prevalence 0.2865   0.1922   0.1759   0.1624   0.1830
## Balanced Accuracy     0.9983   0.9926   0.9933   0.9911   0.9975
```

```
# Run against 20 testing set provided by Professor Leek.
print(predict(testModel3, newdata=pmlTestingData))
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Conclusion

While a model based on Random Forest is slow and uses more processing time, the accuracy jumps from approximately 50% to over 95%, demonstrating that this is a better predictor.